

```
In [13]: import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, datasets, models
from tensorflow.keras.models import Sequential
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
```

```
In [14]: (train_images, train_labels), (test_images, test_labels) = datasets.mnist.load_data()
train_images = train_images.reshape((60000, 28, 28, 1)) / 255.0
test_images = test_images.reshape((10000, 28, 28, 1)) / 255.0
print("TRAIN IMAGES: ", train_images.shape)
print("TEST IMAGES: ", test_images.shape)
```

TRAIN IMAGES: (60000, 28, 28, 1)
TEST IMAGES: (10000, 28, 28, 1)

```
In [15]: model = Sequential([
    layers.Conv2D(64, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(10, activation='softmax')])
```

```
In [16]: model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
    metrics=['accuracy'])
model.summary()
```


Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_4 (Conv2D)	(None, 26, 26, 64)	640
conv2d_5 (Conv2D)	(None, 26, 26, 32)	18464
<hr/>		
conv2d_4 (Conv2D)	(None, 26, 26, 64)	640
conv2d_5 (Conv2D)	(None, 26, 26, 32)	18464
max_pooling2d_3 (MaxPooling 2D)	(None, 13, 13, 32)	0
conv2d_6 (Conv2D)	(None, 13, 13, 16)	4624
max_pooling2d_4 (MaxPooling 2D)	(None, 6, 6, 16)	0
conv2d_7 (Conv2D)	(None, 6, 6, 64)	9280
max_pooling2d_5 (MaxPooling 2D)	(None, 3, 3, 64)	0
flatten_1 (Flatten)	(None, 576)	0
dense_2 (Dense)	(None, 128)	73856
dense_3 (Dense)	(None, 10)	1290
<hr/>		
Total params: 108,154		
Trainable params: 108,154		
Non-trainable params: 0		

```
In [17]: epochs=10
history=model.fit(
    train_images,train_labels,
    epochs=epochs,
    validation_data=(test_images, test_labels)
)
```


```
Epoch 1/10
1875/1875 [=====] - 58s 31ms/step - loss: 0.1347 - accuracy: 0.9566 - val_loss: 0.0418 - val_accuracy: 0.9860
Epoch 2/10
1875/1875 [=====] - 63s 34ms/step - loss: 0.0463 - accuracy: 0.9858 - val_loss: 0.0358 - val_accuracy: 0.9886
Epoch 3/10
1875/1875 [=====] - 52s 28ms/step - loss: 0.0334 - accuracy: 0.9899 - val_loss: 0.0314 - val_accuracy: 0.9905
Epoch 4/10
1875/1875 [=====] - 51s 27ms/step - loss: 0.0253 - accuracy: 0.9919 - val_loss: 0.0556 - val_accuracy: 0.9823
Epoch 5/10
1875/1875 [=====] - 51s 27ms/step - loss: 0.0216 - accuracy: 0.9932 - val_loss: 0.0275 - val_accuracy: 0.9913
Epoch 6/10
1875/1875 [=====] - 51s 27ms/step - loss: 0.0184 - accuracy: 0.9942 - val_loss: 0.0281 - val_accuracy: 0.9910
Epoch 7/10
1875/1875 [=====] - 52s 28ms/step - loss: 0.0160 - accuracy: 0.9948 - val_loss: 0.0270 - val_accuracy: 0.9913
Epoch 8/10
1875/1875 [=====] - 51s 27ms/step - loss: 0.0134 - accuracy: 0.9957 - val_loss: 0.0332 - val_accuracy: 0.9919
Epoch 9/10
1875/1875 [=====] - 51s 27ms/step - loss: 0.0121 - accuracy: 0.9960 - val_loss: 0.0324 - val_accuracy: 0.9900
Epoch 10/10
1875/1875 [=====] - 51s 27ms/step - loss: 0.0098 - accuracy: 0.9968 - val_loss: 0.0314 - val_accuracy: 0.9935
```

```
In [18]: acc=history.history['accuracy']
val_acc=history.history['val_accuracy']
loss=history.history['loss']
val_loss=history.history['val_loss']
epochs_range=range(epochs)
plt.figure(figsize=(8,8))
plt.plot(epochs_range,acc,label='Training Accuracy')
plt.plot(epochs_range,val_acc,label='Validation Accuracy')
plt.plot(epochs_range,loss,label='Training Loss')
plt.plot(epochs_range,val_loss,label='Validation Loss')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy/Loss')
plt.show()
```



```
In [19]: image=train_images[1].reshape(1,28,28,1)
prediction=np.argmax(model.predict(image),axis=-1)
plt.imshow(image.reshape(28,28),cmap='gray')
print('Prediction of model:',prediction[0])
```

```
1/1 [=====] - 0s 69ms/step
Prediction of model: 0
```



```
In [20]: images=test_images[1:5]
for i,test_image in enumerate(images,start=1):
    prediction=np.argmax(model.predict(test_image.reshape(1,28,28,1)),axis=-1)
    plt.subplot(220+i)
    plt.axis('off')
    plt.title("Predicted: {}".format(prediction[0]))
    plt.imshow(test_image.reshape(28, 28),cmap='gray')
    plt.show()
```


1/1 [=====] - 0s 21ms/step

Predicted: 2




1/1 [=====] - 0s 21ms/step

Predicted: 1




1/1 [=====] - 0s 17ms/step

Predicted: 0



1/1 [=====] - 0s 18ms/step


Predicted: 4



```
In [21]: model.save("mnist_cnn.h5")
loaded_model=models.load_model("mnist_cnn.h5")
image=train_images[2].reshape(1,28,28,1)
prediction=np.argmax(loaded_model.predict(image),axis=-1)
plt.imshow(image.reshape(28,28),cmap='gray')
print('Prediction of loaded model:',prediction[0])
```

1/1 [=====] - 0s 75ms/step

Prediction of loaded model: 4



```
In [22]: test_loss,test_acc=model.evaluate(test_images,test_labels,verbose=2)
print("\nTest accuracy:",test_acc)
```


313/313 - 3s - loss: 0.0314 - accuracy: 0.9935 - 3s/epoch - 8ms/step

Test accuracy: 0.9934999942779541

Test accuracy: 0.9934999942779541

```
In [23]: y_pred=np.argmax(model.predict(test_images),axis=-1)
cm=confusion_matrix(test_labels,y_pred)
plt.figure(figsize=(10,8))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```

313/313 [=====] - 3s 10ms/step



```
In [24]: print("\nClassification Report:\n")
print(classification_report(test_labels,y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	980
1	0.99	1.00	1.00	1135
2	1.00	1.00	1.00	1032
3	1.00	1.00	1.00	1010
4	0.99	0.99	0.99	982
5	0.99	0.99	0.99	892
6	1.00	0.99	0.99	958
7	1.00	0.99	0.99	1028
8	0.99	0.99	0.99	974
9	0.99	0.99	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000